

# Challenging Task Brief: Project Vaani Sentinel X

**Duration:** ~16–18 Hours

**Level:** Advanced

**Track:** AI x Cybersecurity x Full Stack Automation x Voice Interfaces

## Project Context for Candidate (Visible)

You're tasked with building an AI-integrated voice-first platform that auto-generates platform-specific content from structured knowledge, simulates social media scheduling, and ensures end-to-end modularity and security. This is meant to serve as a blueprint for deploying secure AI systems in real-time public environments.

This is not a UI-centric project. Your focus is backend intelligence, agent orchestration, and voice-driven interaction pipelines. You're encouraged to be modular, agent-driven, secure, and scalable in your design.

## Your Mission – Build an Autonomous Agent-Based System with the Following Flow:

### 1. Agent A: Knowledge Miner & Sanitizer

- Accept raw CSV or JSON data input (facts, quotes, micro-articles, etc.)
- Structure it into modular “content blocks” for further processing
- Add a content verification layer:
  - Auto-check for profanity
  - Detect if content is neutral/biased
  - Validate against a dummy “truth-source.csv” file

### 2. Agent B: AI Writer & Voice Synth Generator

- Use OpenAI API (or a local model if offline) to:
  - Convert each block into:
    - A tweet ( $\leq 280$  characters)
    - A 1-paragraph post for Instagram/LinkedIn

- A 20–30 second voice script for use in an AI assistant
- Generate the TTS (Text-to-Speech) file using:
  - Vapi / IBM Watson / Google TTS / ElevenLabs (if access)
- Save outputs in a /content\_ready folder, maintaining versioning

### **3. Agent C: Secure Web Interface (Minimalist UI)**

- Next.js / React frontend with:
  - Login system (Firebase or JWT)
  - Panel to view and play AI-generated content (with audio)
  - Download / Copy-ready buttons for each content type
  - Show publishing time and versioning

### **4. Agent D: Scheduler & Publisher Simulator**

- Schedule each post as if it would be published on:
  - Twitter (tweet)
  - Instagram (longer post)
  - Spotify/Voice Platforms (TTS audio)
- Simulate these by:
  - Storing to scheduled\_posts/
  - Creating dummy POST calls to mock endpoints
  - Logging time, platform, and result to Firebase or a local DB

### **5. Agent E: Security & Ethics Guard (Advanced Subsystem)**

- Add a mini-AI/regex tool to:
  - Flag controversial content (religion, politics, bias)
  - Simulate an alert dashboard (even if backend-only)
  - Encrypt content archives with basic Python encryption
  - Include a “kill switch” to wipe data if misuse is detected
  -

## Stack Expectations

- Frontend: React.js / Next.js (minimal UI)
- Backend: Python (FastAPI), Firebase OR local storage
- Voice: IBM Watson TTS / Web Speech API / Vapi
- AI: OpenAI, Ollama, or local LLMs (text generation)
- Security: Regex + encryption + flagging heuristics
- DevOps: GitHub repo + README + folder structure

## Bonus Challenges (Optional, Worth Extra Credit)

- Add logging system with user ID, timestamps, and agent actions
- Create your own “Command Center CLI” to run agents individually
- Create a dashboard showing how each post scored on “ethics”, “virality”, and “neutrality”

## Folder Structure Expected

```
vaani-sentinel-x/  
├── agents/  
│   ├── miner_sanitizer.py  
│   ├── ai_writer_voicegen.py  
│   ├── scheduler.py  
│   ├── publisher_sim.py  
│   └── security_guard.py  
├── web-ui/  
│   └── nextjs-voice-panel/  
├── content/  
│   ├── raw/  
│   ├── structured/  
│   └── content_ready/  
├── logs/  
├── scheduler_db/  
├── kill_switch.py  
└── README.md
```

# Deliverables

- GitHub repo with full code + setup steps
- A video demo (screen recording, 3–5 min)
- Report/README outlining how agents work, what libraries were used, and any blockers you faced

# Evaluation Rubric

Category	Points
Modular Agent Design	20
AI + TTS Integration	20
Secure Content Processing	20
Voice-Ready UI	15
Automation & Scheduling	15
Bonus Implementations	10
Total	100